

MLW / KUHeS Statistics and R short course

Session 1 - Practical (solutions)

Evaristar Kudowa, Marc Henrion

2026-01-19

Session 1 - Practical (Solutions)

Go to the course website on GitHub:

https://github.com/mlw-stats/R_And_Statistics_Training_2026/Session1

From here, download the following files:

`btTBreg.csv`
`btTBregHospitals.csv`
`btTBreg_info.txt`

1. Load the `btTBreg.csv` data table into R.

```
btDat<-read.csv("dataAndSupportDocs/btTBreg.csv")
```

```
head(btDat) # have a look at the data
```

	id	age	sex	hiv	bmi	ses	cd41	cd42	cd41.sk	cd42.sk	hosp
1	1	44	2	0	26.32	4	346	519	313.11656	572.8906	1
2	2	32	2	0	20.79	5	237	337	43.12752	406.1971	5
3	3	32	1	0	19.21	1	198	328	338.32172	408.2427	2
4	4	20	1	0	21.34	4	246	525	77.08697	312.7572	3
5	5	30	1	0	23.98	4	270	444	169.02539	335.3739	3
6	6	32	1	0	17.97	4	283	372	255.45773	323.4773	4

```
dim(btDat) # check dimesnions of data table
```

```
[1] 3000    11
```

2. The variables `cd41`, `cd42` and `cd41.sk`, `cd42.sk` measure the same variables (`cd4` and `cd4.sk` respectively) in the same individuals at two different time point. This means the data are in wide format. Reformat to long format.

The key difficulty here is that you have 2 variables (at 2 times points). In the example from lectures we only had 1 variable (at 3 different conditions). One approach is to do each variable separately, then combine the resulting data frames:

```
btDatLong.cd4<-btDat %>%
  pivot_longer(names_to="time", values_to="cd4", cols=c(cd41, cd42)) %>%
  select(id,age,sex,hiv,bmi,ses,hosp,time,cd4)

btDatLong.cd4sk<-btDat%>%
  pivot_longer(names_to="time", values_to="cd4.sk", cols=c(cd41.sk, cd42.sk)) %>%
  select(id,age,sex,hiv,bmi,ses,hosp,time,cd4.sk)

btDatLong<-data.frame(btDatLong.cd4,cd4.sk=btDatLong.cd4sk$cd4.sk)

rm(btDatLong.cd4,btDatLong.cd4sk)

btDatLong$time<-factor(
  case_when(
    btDatLong$time=="cd41"~"entry",
    btDatLong$time=="cd42"~"exit",
    TRUE~NA_character_)
  ) # rename the levels of the time variable

head(btDatLong) # have a look at the data
```

	id	age	sex	hiv	bmi	ses	hosp	time	cd4	cd4.sk
1	1	44	2	0	26.32	4	1	entry	346	313.11656
2	1	44	2	0	26.32	4	1	exit	519	572.89062
3	2	32	2	0	20.79	5	5	entry	237	43.12752
4	2	32	2	0	20.79	5	5	exit	337	406.19707
5	3	32	1	0	19.21	1	2	entry	198	338.32172
6	3	32	1	0	19.21	1	2	exit	328	408.24267

```
dim(btDatLong) # check dimensions
```

```
[1] 6000 10
```

This can be done a bit more directly, by using regular expression (character expressions that match flexibly to names) and a combination of `pivot_longer()` and `pivot_wider()`:

```
btDatLong<-btDat %>%
  pivot_longer(cols=c(cd41, cd42, cd41.sk, cd42.sk),
               names_pattern = "cd4(1|2)(.*)",
               names_to = c("time", "cd4"),
               values_to="measurement") %>%
  mutate(cd4=paste(sep="", "cd4", cd4)) %>%
  pivot_wider(names_from=cd4, values_from=measurement)
```

```
head(btDatLong) # have a look at the data
```

```
# A tibble: 6 x 10
  id    age   sex   hiv   bmi   ses hosp time     cd4 cd4.sk
  <int> <int> <int> <int> <dbl> <int> <int> <chr> <dbl> <dbl>
1    1     44     2     0  26.3     4     1  1      346  313.
2    1     44     2     0  26.3     4     1  2      519  573.
3    2     32     2     0  20.8     5     5  1      237  43.1
4    2     32     2     0  20.8     5     5  2      337  406.
5    3     32     1     0  19.2     1     2  1      198  338.
6    3     32     1     0  19.2     1     2  2      328  408.
```

```
dim(btDatLong) # check dimensions
```

```
[1] 6000    10
```

The code above requires a bit of unpacking:

- The expression in brackets in the “names_pattern” argument are regular expression matching sequences of character: “(1|2)” matches 1 or 2 and “(.)” matches anything.
- The “`mutate()`” line is needed as the values stored in the “cd4” column are “” and “.sk” – as the `pivot_wider()` statement on the next line will use those as column names, we cannot have an empty column name – “” would trigger an error message. So we just add the characters “cd4” in front of the stored values – i.e. we then have “cd4” and “cd4.sk” rather than “” and “.sk”.

An alternative function that can be used is `reshape()`. To get more information on this function, type `?reshape` at the console.

```

btDatLong<-reshape(btDat,
                     direction="long",
                     varying=list(c("cd41","cd42"),c("cd41.sk","cd42.sk")),
                     id=id,
                     v.names=c("cd4","cd4.sk"))

head(btDatLong) # have a look at the data

```

	id	age	sex	hiv	bmi	ses	hosp	time	cd4	cd4.sk
1.1	1	44	2	0	26.32	4	1	1	346	313.11656
2.1	2	32	2	0	20.79	5	5	1	237	43.12752
3.1	3	32	1	0	19.21	1	2	1	198	338.32172
4.1	4	20	1	0	21.34	4	3	1	246	77.08697
5.1	5	30	1	0	23.98	4	3	1	270	169.02539
6.1	6	32	1	0	17.97	4	4	1	283	255.45773

```
dim(btDatLong) # check dimensions
```

[1] 6000 10

3. Save the reformatted data into a file called `btTBregLong.tab` in such a way that
 - i. Columns are tab-separated.
 - ii. Column names are saved.
 - iii. No row number is saved in the resulting file.

```

dir.create("Session1_output", showWarnings=F)
write.table(btDatLong, sep="\t", col.names=T, row.names=F, file="Session1_output/btTBregLong.tab")

```

4. Copy the code below to generate some wide-format data. We will assume this dataset contains observations of 2 biomarkers, `ferritin` and `rbp4` for 10 study participants at 2 different timepoints, `day1` and `day90`.

```

set.seed(123)

df<-data.frame(
  id=paste(sep="", "P", 1:10),
  ferritin_day1=rexp(10, rate=1/195),
  rbp4_day1=rexp(10, rate=1/2.5)
) %>%
  mutate(

```

```

    ferritin_day90=rnorm(10,mean=ferritin_day1+5,sd=4),
    rbp4_day90=rbp4_day1+rexp(10,rate=1/0.25)
)

```

This is what this data table looks like:

	id	ferritin_day1	rbp4_day1	ferritin_day90	rbp4_day90
1	P1	164.474166	2.5120751	169.251422	3.021712
2	P2	112.439003	1.2005368	114.301474	1.552174
3	P3	259.165699	0.7025341	261.231686	0.829188
4	P4	6.157585	0.9427946	10.294123	1.007684
5	P5	10.961140	0.4707101	14.621489	1.119933
6	P6	61.717737	2.1244653	62.374941	2.431722
7	P7	61.274322	3.9080088	65.932629	4.105679
8	P8	28.327027	1.1969010	37.609469	1.354221
9	P9	531.616111	1.4773371	536.034536	1.790997
10	P10	5.684922	10.1025293	6.022743	10.249700

Reformat this to long format, i.e. so that you have 4 columns: **id**, **time**, **ferritin** and **rbp4**.

```

dfLong<-df %>%
  pivot_longer(cols=c(ferritin_day1,rbp4_day1,ferritin_day90,rbp4_day90),
               names_pattern="(.*)_(.*)",
               names_to=c("biomarker","timepoint"),
               values_to="value") %>%
  pivot_wider(names_from=biomarker,values_from=value)

print(dfLong)

```

	# A tibble: 20 x 4	id	timepoint	ferritin	rbp4
		<chr>	<chr>	<dbl>	<dbl>
1	P1	day1		164.	2.51
2	P1	day90		169.	3.02
3	P2	day1		112.	1.20
4	P2	day90		114.	1.55
5	P3	day1		259.	0.703
6	P3	day90		261.	0.829
7	P4	day1		6.16	0.943
8	P4	day90		10.3	1.01
9	P5	day1		11.0	0.471

```

10 P5      day90        14.6   1.12
11 P6      day1         61.7   2.12
12 P6      day90        62.4   2.43
13 P7      day1         61.3   3.91
14 P7      day90        65.9   4.11
15 P8      day1         28.3   1.20
16 P8      day90        37.6   1.35
17 P9      day1         532.   1.48
18 P9      day90        536.   1.79
19 P10     day1         5.68  10.1
20 P10     day90        6.02  10.2

```

5. Load the `btTBregHospitals.csv` data table. Join the data frames storing `btTBreg.csv` and `btTBregHospitals.csv`.

```
btDatHosp<-read.csv("dataAndSupportDocs/btTBregHospitals.csv")
```

```
head(btDatHosp) # have a look at the data
```

	HID	ShortName	FullName	beds	city
1	1	QECH	Queen Elizabeth Central Hospital	1000	Blantyre
2	2	KCH	Kamuzu Central Hospital	1000	Lilongwe
3	3	ZCH	Zomba Central Hospital	400	Zomba
4	4	MCH	Mzuzu Central Hospital	350	Mzuzu
5	5	Mlambe	Mlambe Mission Hospital	254	Lunzu

```
dim(btDatHosp) # check dimensions of the data table
```

```
[1] 5 5
```

```
btDatJoined<-btDat %>%
  inner_join(btDatHosp,by=c("hosp"="HID"))
```

```
head(btDatJoined) # have a look
```

	id	age	sex	hiv	bmi	ses	cd41	cd42	cd41.sk	cd42.sk	hosp	ShortName
1	1	44	2	0	26.32	4	346	519	313.11656	572.8906	1	QECH
2	2	32	2	0	20.79	5	237	337	43.12752	406.1971	5	Mlambe
3	3	32	1	0	19.21	1	198	328	338.32172	408.2427	2	KCH
4	4	20	1	0	21.34	4	246	525	77.08697	312.7572	3	ZCH

```

5 5 30 1 0 23.98 4 270 444 169.02539 335.3739 3 ZCH
6 6 32 1 0 17.97 4 283 372 255.45773 323.4773 4 MCH
          FullName beds city
1 Queen Elizabeth Central Hospital 1000 Blantyre
2 Mlambe Mission Hospital 254 Lunzu
3 Kamuzu Central Hospital 1000 Lilongwe
4 Zomba Central Hospital 400 Zomba
5 Zomba Central Hospital 400 Zomba
6 Mzuzu Central Hospital 350 Mzuzu

```

```
dim(btDatJoined) # check dimensions
```

```
[1] 3000 15
```

- Compute the average patient age and the proportion of male patients for each hospital.

Useful functions for this are `aggregate()` and `group_by()`. You can however also do it manually.

- Manually:

```

# initialise new variables
btDatHosp$avgAge<-NA
btDatHosp$propMale<-NA

# iterate over hospitals
for(i in 1:nrow(btDatHosp)){
  btDatHosp$avgAge[i]<-mean(btDatJoined$age[btDatJoined$ShortName==btDatHosp$ShortName[i]],na.rm=TRUE)
  btDatHosp$propMale[i]<-sum(btDatJoined$sex==1 &
                                btDatJoined$ShortName==btDatHosp$ShortName[i]) /
                                sum(btDatJoined$ShortName==btDatHosp$ShortName[i])
}

print(btDatHosp)

```

	HID	ShortName	FullName	beds	city	avgAge
1	1	QECH	Queen Elizabeth Central Hospital	1000	Blantyre	33.14020
2	2	KCH	Kamuzu Central Hospital	1000	Lilongwe	32.80067
3	3	ZCH	Zomba Central Hospital	400	Zomba	32.99310
4	4	MCH	Mzuzu Central Hospital	350	Mzuzu	32.87382
5	5	Mlambe	Mlambe Mission Hospital	254	Lunzu	32.89950
		propMale				

```
1 0.4763514
2 0.4757119
3 0.4948276
4 0.4731861
5 0.5242881
```

- Using `aggregate()`

```
btDat$hosp<-factor(btDat$hosp)
btDatHosp$avgAge<-aggregate(btDatJoined$age,FUN=mean,by=list(btDat$hosp))$x
btDatHosp$propMale<-aggregate(ifelse(btDatJoined$sex==1,1,0),FUN=mean,by=list(btDat$hosp))$x

print(btDatHosp)
```

HID	ShortName	FullName	beds	city	avgAge
1	1	QECH Queen Elizabeth Central Hospital	1000	Blantyre	33.14020
2	2	KCH	1000	Lilongwe	32.80067
3	3	ZCH	400	Zomba	32.99310
4	4	MCH	350	Mzuzu	32.87382
5	5	Mlambe	254	Lunzu	32.89950

```
propMale
1 0.4763514
2 0.4757119
3 0.4948276
4 0.4731861
5 0.5242881
```

- Using `group_by()`

```
tmp<-btDat %>%
  group_by(hosp) %>%
  summarise(avgAge=mean(age,na.rm=T))
btDatHosp$avgAge<-tmp$avgAge

tmp<-btDat %>%
  group_by(hosp) %>%
  summarise(propMale=mean(ifelse(sex==1,1,0),na.rm=T))
btDatHosp$propMale<-tmp$propMale

print(btDatHosp)
```

	HID	ShortName		FullName	beds	city	avgAge
1	1	QECH	Queen Elizabeth Central Hospital	1000	Blantyre	33.14020	
2	2	KCH	Kamuzu Central Hospital	1000	Lilongwe	32.80067	
3	3	ZCH	Zomba Central Hospital	400	Zomba	32.99310	
4	4	MCH	Mzuzu Central Hospital	350	Mzuzu	32.87382	
5	5	Mlambe	Mlambe Mission Hospital	254	Lunzu	32.89950	
		propMale					
1		0.4763514					
2		0.4757119					
3		0.4948276					
4		0.4731861					
5		0.5242881					

7. Write an R function that computes the following summary statistics, then, using your custom function, compute these for the `bmi`, `cd41`, `cd42` columns:

- mean
- median
- interquartile range
- minimum
- maximum
- number of missing values

```
summaryFun<-function(x){
  return(c(
    mean(x,na.rm=T),
    median(x),
    paste(sep="", "(" ,paste(collapse=",",quantile(x,probs=c(0.25,0.75))),")"),
    min(x,na.rm=T),
    max(x,na.rm=T),
    sum(is.na(x)))
  ))
}

res<-apply(btDat[,c("bmi","cd41","cd42")],MARGIN=2,FUN=summaryFun)
rownames(res)<-c("mean","median","IQR","min","max","num_MV")
print(res)
```

	bmi	cd41	cd42
mean	"23.0574333333333"	"248.79433333333"	"448.003"
median	"23.05"	"249"	"447"
IQR	"(21.34,24.74)"	"(216,281)"	"(381,515)"
min	"12.64"	"57"	"81"

```

max      "31.14"          "447"          "843"
num_MV  "0"              "0"            "0"

```

8. Do the same now, but only for female patients. Repeat for only male patients.

```

resF<-apply(btDat[btDat$sex==2,c("bmi","cd41","cd42")],MARGIN=2,FUN=summaryFun)
rownames(resF)<-c("mean","median","IQR","min","max","num_MV")
print(resF)

```

	bmi	cd41	cd42
mean	"23.1218644067797"	"248.473924380704"	"446.675358539765"
median	"23.14"	"250"	"447.5"
IQR	"(21.365,24.82)"	"(215,281)"	"(379,512)"
min	"12.64"	"57"	"138"
max	"31.14"	"447"	"820"
num_MV	"0"	"0"	"0"

```

resM<-apply(btDat[btDat$sex==1,c("bmi","cd41","cd42")],MARGIN=2,FUN=summaryFun)
rownames(resM)<-c("mean","median","IQR","min","max","num_MV")
print(resM)

```

	bmi	cd41	cd42
mean	"22.9900136425648"	"249.129604365621"	"449.392223738063"
median	"22.98"	"248"	"447"
IQR	"(21.3,24.66)"	"(216,282)"	"(383,519.75)"
min	"14.44"	"71"	"81"
max	"30.9"	"414"	"843"
num_MV	"0"	"0"	"0"